


# Advanced Artificial Intelligence

## CM4107 (Week 3)



Searching and Planning, Computational Game Theory

Dr. Yann Savoye

School of Computing Science and Digital Media

Robert Gordon University

# Module Information

- **Assessment:**

- Coursework (2 components)
  - Component 1: literature review
  - Component 2: paper implementation
- No mid-term or final written exam.

*All deadlines are strong:*

- *It will not be possible to upload material after the deadline.*
- *No deadline extension will be granted. No excuse.*
- *Only the content submitted via the Moodle will be mark.*

# Coursework

- **Submission of the Coursework Part 1:**

**Deadline - Monday, October 30th, 2017 23:00:**

*Activity 1 and Activity 2*

- *2-pages written reports (in PDF format)*
- *7 slides presentation (in PDF format)*


- You should have selected two papers! And started to read them ...

# Coursework

- How to read a paper? **“Three Passes Technique”**



- *First Pass : problem, key ideas*



- *Second Pass : high-level pipeline*




- *Third Pass : low-level details*

# Overview

- Part I – Computational Game Theory
- Part II – Game Formalization and Solving Process
- Part III – Motion Planning and Behaviours
- Part IV – Searching
- Part V – Optimizing
- Part VI – Conclusions


# Part I – Computational Game Theory

# Theory of Games



***Game Theory*** was developed by ***John Von Neumann*** in 1944


# Nash Equilibrium



*One of the fundamental principle of **game theory** , the idea of **equilibrium** was developed by **John Nash***



# Games



*Beyond what we call '**games**', such as chess, poker, soccer, it includes the **modelling of conflict***

# What is a Game?

- A game : an **abstract model** for interacting **self-interested** players.
- **Abstract** = model **relevant to the decisions** that players make
- Focus on **decision-making** influencing the **outcomes**
- **Self-interest**: Players assumed to **act in their own interests**
- **Self-interested agents** : the game models an **optimization process**.
- A **game** consist of a set of **players, strategies, payoff**.
- **Players**: everyone who has an effect on your earnings
- **Strategies**: actions available to each player
- **Payoffs**: **Reflect the interests/outcome of the players**
- **Cooperative Game**: players form binding commitments
- **Non-cooperative Game**: player make decision independently.
- **Bayesian games, repeated and stochastic games**

# What is game theory?



***Game Theory provides some interesting lessons in strategic thinking.***

# What is game theory?

- **Game Theory = the theory of rational decision in conflicts.**
- Mathematics of determining **strategies** for optimal play.
- **Strategies for dealing with competitive situations.**
- Anticipation of the adversary's play and respond accordingly.
- **AI community = dominant formalism for studying strategic.**
- **Game theory in various research discipline.**
- **Game-theoretic mechanisms in Economics.**
- **Computational Game Theory: address algorithmic issues.**
- **Optimize an objective function.**
- **Military battles, business interactions, managerial economics**

# Brief introduction to Game theory

- **Game theory is the mathematical modelling of strategy**
- **Decision making**
- Popularized by **movies** such as "**A Beautiful Mind**".
- **Interaction between agents** to achieve good outcomes
- **Strategic interaction** among rational (and irrational) agents.
- **Social choice theory** (i.e., collective decision making)
- **Representing game and strategy**
- The **optimal strategy** is based on the **minimax concept**
- Each player **maximizes the minimum** values obtainable.
- Designing **mechanisms to maximize** aggregate happiness.

# Key Concepts

Cooperative outcome	An equilibrium in a game where the players agree to cooperate
Dominant strategy	A dominant strategy is one where a single strategy is best for a player regardless of what strategy other players in the game decide to use
Nash equilibrium	Any situation where all participants in a game are pursuing their best possible strategy given the strategies of <u>all</u> of the other participants
Tacit collusion	Where firms undertake actions that are likely to minimize a competitive response, e.g. avoiding price-cutting or not attacking each other's market
Whistle blowing	When one or more agents in a collusive agreement report it to the authorities
Zero sum game	An economic transaction in which whatever is gained by one party must be lost by the other.

# Applications of Game Theory

- **Economy**
- Politics (vote, coalitions)
- Biology (Darwin's principle, evolutionary)
- **Anthropology**
- **War**
- **Management-labor arbitration**
- **Philosophy (morality and free-will)**
- **National Football league**
- **Trading**



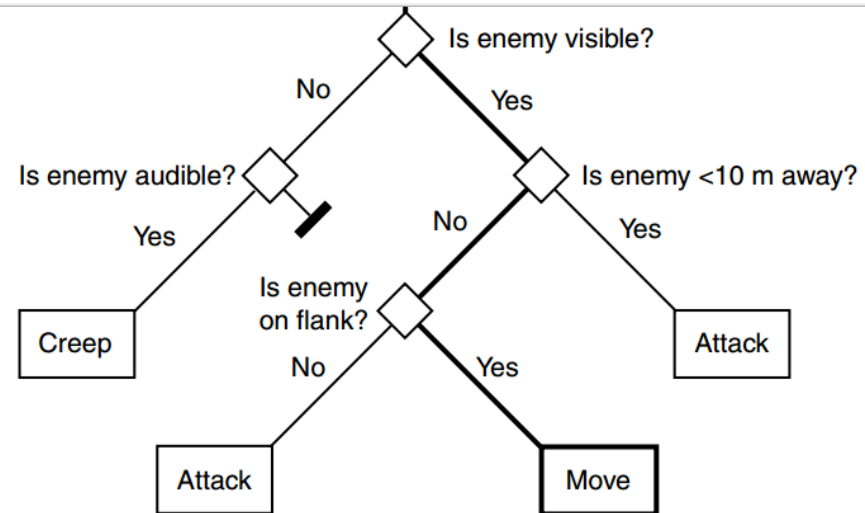
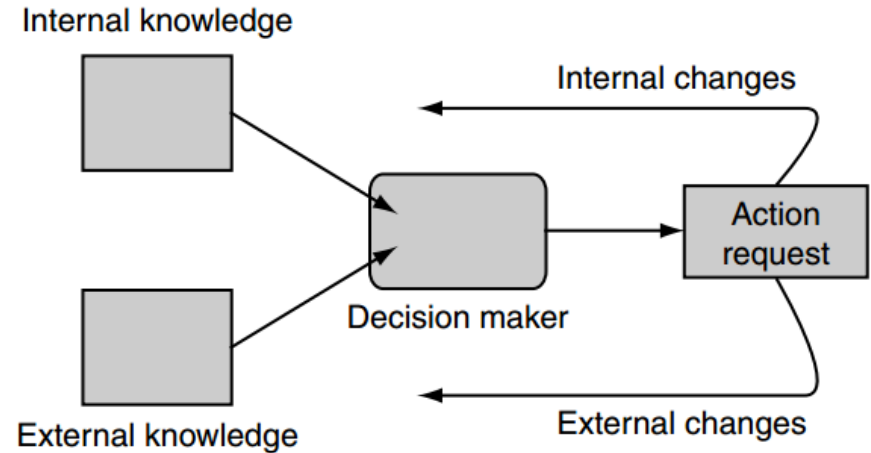
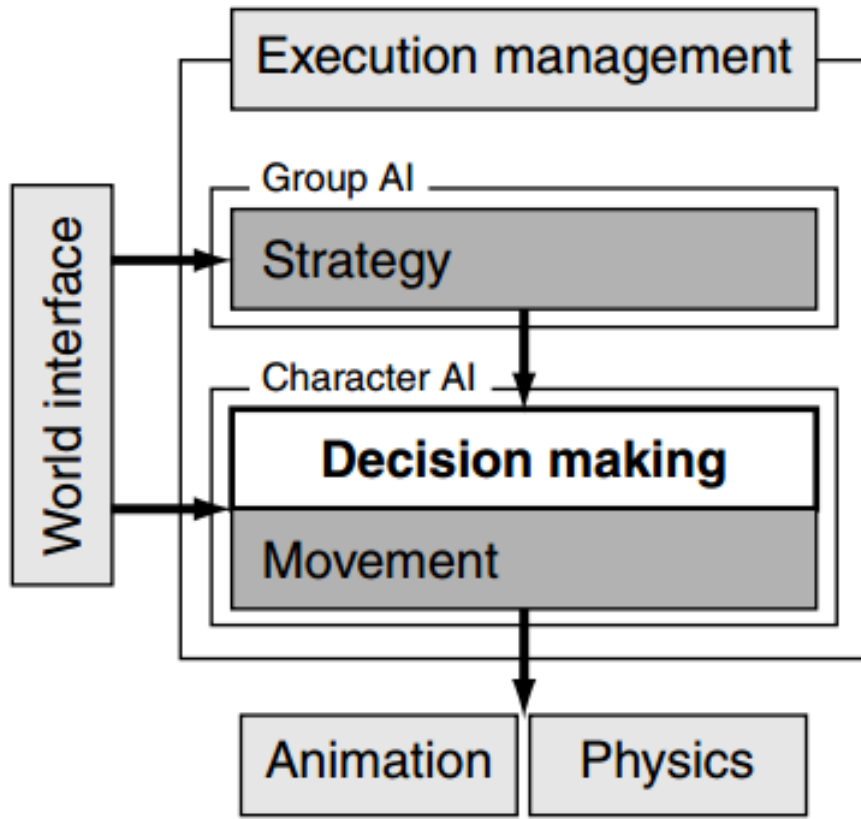
***Gamification is the process of using game-based mechanism to solve problem***

# Strategies in Game

- **Kakutani's** fixed point / **Minimax** / **Arrow** theorem
  - **Dominance** solvability
  - **Rationalizability**
  - **Social choice** theory
  - **Evolutionary** stable states and sets
  - **Nash equilibrium**
- 
- **Strategies are evaluated along the following dimensions**
    - Completeness**: always find a solution if exists?*
    - Time complexity**: time a solution*
    - Space Complexity**: enough memory ?*
    - Optimality** : is the solution improvable ?*



# Decision Making



# Importance of Game Theory in AI

- **Game theory** and AI is a way to think and **model systems**.
- **Games** = the most **areas of progress** in the **AI space**
- Lot of **games mastered by AI** systems using break through algos.
- The **success of AI** is really tied to the progress on **game theory**.
- **Games** are a **materialization of game theory**.
- **Collaborating** or **competing** to accomplish a task can be gamified
- Several **aspects of game theory** that could help to **understand AI**
- Helps agents **select strategies**
- Guarantees about **artificially designed mechanisms**
- Automated analysis of **strategic models**

# Dominance and Payoff Matrix

- **Dominance** = when one strategy is better than another strategy
- A strategy is dominant if a player earn a **larger payoff** than any other.
- A **payoff matrix** describes the **outcome** for each player and for each **set of strategic choices**.



# The 3-doors Monty Hall Problem



*You're given the choice of three doors: Behind one door is a **car**; behind the others, **goats**.*


*Is it to your advantage to switch your choice?*

	DOOR 1	DOOR 2	DOOR 3	RESULT
GAME 1	Auto	Goat	Goat	Switch → <b>LOSE</b>
GAME 2	Goat	Auto	Goat	Switch → <b>WIN</b>
GAME 3	Goat	Goat	Auto	Switch → <b>WIN</b>
GAME 4	Auto	Goat	Goat	Stay → <b>WIN</b>
GAME 5	Goat	Auto	Goat	Stay → <b>LOSE</b>
GAME 6	Goat	Goat	Auto	Stay → <b>LOSE</b>

*the dominance implies that a strategy maximizing the probability of winning the car.*

**always-switching strategies**

# Rock Paper Scissors



Gains of Player 1

Player 2 chooses...



Player 1 chooses...



0, 0	1, -1	1, -1
-1, 1	0, 0	-1, 1
-1, 1	1, -1	0, 0

# Prisoners' Dilemma



Betray




Cooperate

		Prisoner B	
		Confess	Keep quiet
Prisoner A	Confess	Both go to jail for ten years	Prisoner B gets life imprisonment, A goes free
	Keep quiet	Prisoner A gets life imprisonment, B goes free	Both go to jail for one year

*The **prisoner's dilemma** is a paradox in which two individuals acting in their own **self-interest** pursue a course of action that does not result in the **ideal outcome***

# Nash Equilibrium



## *EQUILIBRIUM POINTS IN N-PERSON GAMES*

BY JOHN F. NASH, JR.\*

PRINCETON UNIVERSITY

Communicated by S. Lefschetz, November 16, 1949

One may define a concept of an  $n$ -person game in which each player has a finite set of pure strategies and in which a definite set of payments to the  $n$  players corresponds to each  $n$ -tuple of pure strategies, one strategy being taken for each player. For mixed strategies, which are probability

***Nash Theorem : every game with a finite number of players and strategies .  
has at least one equilibrium (in pure or mixed strategy).***

- **John Nash** = proved the existence of an equilibrium point in non-cooperative games
- The defense adjust its pressure until the efficiencies of the offensive are equal

## Part II – Game Formalization and Solving Process



# Formulating the 8-Puzzle Problem

- **8-puzzle** is a **sliding puzzle** that consists of a frame of **numbered square tiles in random order** with one tile missing.
- **States:** each represented by a 3x3 array of numbers in  $[0...8]$ , where value 0 is for the empty cell.




$$A = \begin{matrix} 4 & 1 & 2 \\ 0 & 8 & 7 \\ 6 & 3 & 5 \end{matrix}$$

# Problem Space Model


A problem is defined by 5 items:

- Space of states
- **Initial state** (given)
- **Goal State** (given)
- Transition model
- Path

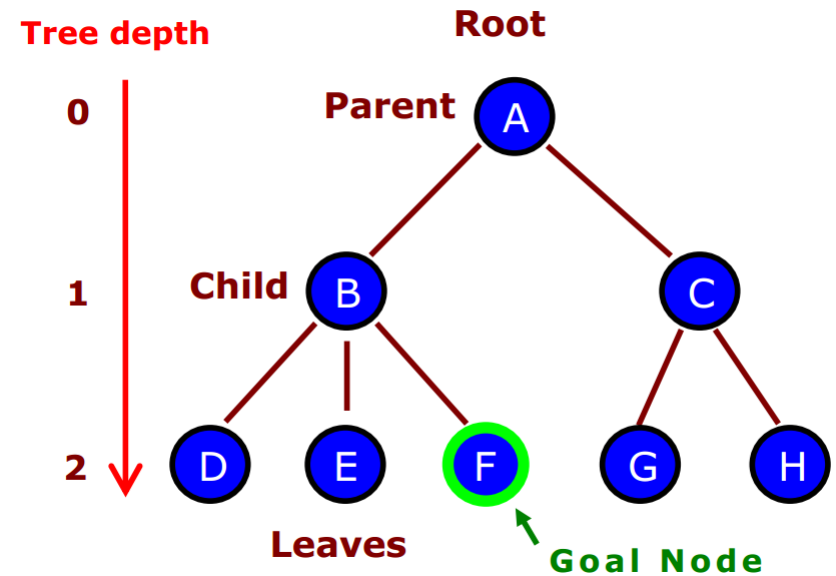


*A **solution** is a sequence of actions leading from the initial state to a goal state.*

# Representation of States




**Graph**



**Tree**


# Adversarial Games



(image courtesy of Zoran Popovic)


- **Exploit** the possible **strategies of an adversary**
- Controllers for characters in **competitive games**.
- Results **goals are in conflict**, giving rise to **adversarial search problem**
- Finds the move **that minimizes the maximum** utility of an **opponent**.

# Alpha-Beta Pruning



- **Minimax:** to **minimize** the maximum **loss** (defensive)
- **Maximin:** to **maximize** the minimum **gain** (offensive)
- Minimax = Maximin
- Minimax search = **exponential in depth** of the tree.
- **Pruning:** Eliminate some part of the tree.

# Complexity of Solving Games



*Optimal solution of  $n$ -Puzzle family is **NP-hard**.*

# Complexity of Solving Games

- **Computational complexity of solving games**
- Determine if **game theory** can be used to **model real-world** settings
- **Solving** requires the use of **computational power**
- **Solution concept is realistic**
- **Complexity of solving** gives a **lower bound on complexity**

$$\textit{Algorithmic Game Theory} = \textit{Theory of Algorithms} + \textit{Game Theory}$$

# Solution Concepts



Goal  
Satisfaction

reach the goal node  
Constraint satisfaction

Optimization


optimize(objective fn)  
Constraint Optimization

- **Existence:** Does a solution guarantee that the game is rational?
- **Uniqueness:** Does the solution is unique?
- **Tractability:** Is it computationally easy to verify?
- **Comprehensibility:** Is it easy to understand?
- **Invariance:** Do a small change in the setup lead to different solution?
- **Search and Optimization**



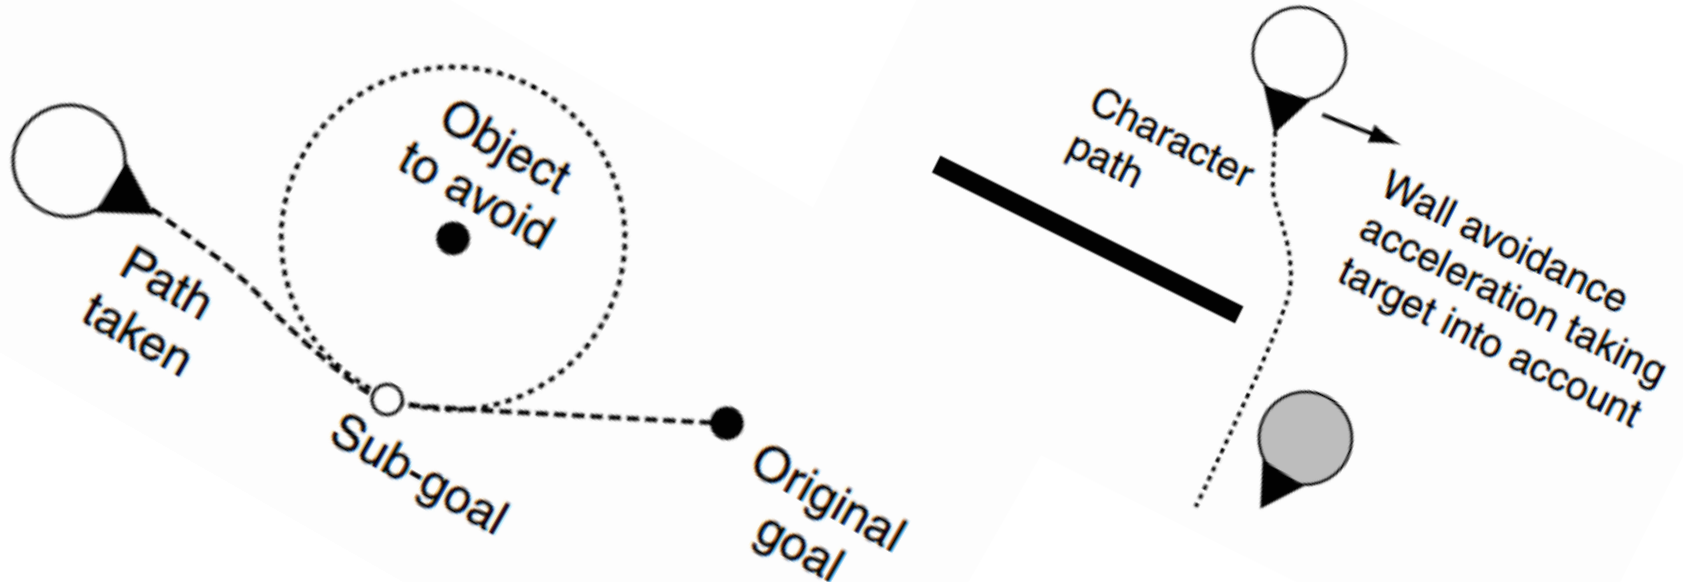
## Part III – Motion Planning and Behaviors

# Motion Planning





- **Motion planning** = breaking down a desired movement task
- Satisfy **movement constraints** and **optimize** the movement.
- **Motion planning** = algorithm about the movement.
- Development of **motion planning algorithms** for **physical robots**

# Collision and Wall Avoidance



In robotics, **obstacle avoidance** is the task of satisfying some control objective subject to non-intersection or non-collision position constraints.

# Flocking/ Steering Behaviors



- **Drives the flock** in the direction of a target/goal.
- Rules of Flocking Behaviors: **Alignment, Cohesion, and Separation**
- **Steering behaviours** control **goal-directed motions**.


# A\* Path Finding Algorithm

## Algorithm 24 A\* Algorithm

**Input:** A graph

**Output:** A path between start and goal nodes

```
1: repeat
2:   Pick  $n_{best}$  from  $O$  such that  $f(n_{best}) \leq f(n), \forall n \in O$ .
3:   Remove  $n_{best}$  from  $O$  and add to  $C$ .
4:   If  $n_{best} = q_{goal}$ , EXIT.
5:   Expand  $n_{best}$ : for all  $x \in \text{Star}(n_{best})$  that are not in  $C$ .
6:     if  $x \notin O$  then
7:       add  $x$  to  $O$ .
8:     else if  $g(n_{best}) + c(n_{best}, x) < g(x)$  then
9:       update  $x$ 's backpointer to point to  $n_{best}$ 
10:    end if
11: until  $O$  is empty
```




- Finding **The Shortest Path**
- **Problem of finding a path** between two vertices (or nodes) in a graph
- **Cost function:** the sum of edges' weights is minimized.



## Part IV – Searching

# Binary Search Tree

```
BinarySearch(list[], min, max, key)
if max < min then
    return false
else
    mid = (max+min) / 2
    if list[mid] > key then
        return BinarySearch(list[], min, mid-1, key)
    else if list[mid] < key then
        return BinarySearch(list[], mid+1, max, key)
    else
        return mid
    end if
end if
```




# Backtracking Search Algorithm



***Backtracking search** is used for a **depth-first search** that chooses values for one variable at a time and **backtracks** when a variable has no legal values left to assign.*





# Line Search Methods




$$\min_{\alpha > 0} f(x_k + \alpha p_k).$$

***Line search method*** seeks the minimum of a defined nonlinear function by selecting a ***reasonable direction vector***

# Local vs Global Maxima/Minima



**Can get stuck in a local maximum**



**Can get stuck on a plateau**

# Greedy Search

```
Algorithm Greedy (a,n)
//a[1:n]contains the n inputs.
{
    solution:=0;//initialize the solution.
    for i:=1 to n do
    {
        x:=Select(a);
        if Feasible( solution, x) then
            solution:=Union(solution,x);
    }
    return solution;
}
```

A **greedy** search follows a **heuristic** by making the **locally optimal** choice at each stage with the hope of **finding a global optimum**

# Iterated Local Search

**procedure** *Iterated Local Search*

$s_0 \leftarrow \text{GenerateInitialSolution}$

$s^* \leftarrow \text{LocalSearch}(s_0)$

**repeat**


$s' \leftarrow \text{Perturbation}(s^*, \text{history})$

$s^{*'} \leftarrow \text{LocalSearch}(s')$

$s^* \leftarrow \text{AcceptanceCriterion}(s^*, s^{*'}, \text{history})$

**until** termination condition met

**end**

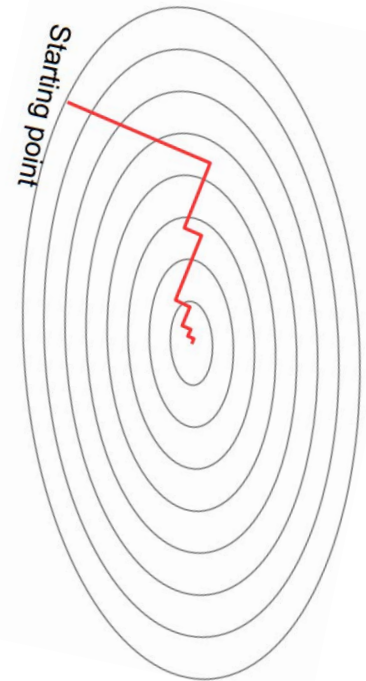


- **Iterated Local Search** = a sequence of **locally optimal** solutions
- The **modified solution** = **Perturbing** the current local minimum
- The perturbation strength has to be sufficient to **lead the trajectory**

# Hill Climbing

```
function HILL-CLIMBING(problem) returns a solution state
  inputs: problem, a problem
  static: current, a node
           next, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    next ← a highest-valued successor of current
    if VALUE[next] < VALUE[current] then return current
    current ← next
  end
```




- **Hill climbing** is a **local search technique**.
- Hill Climbing **gets stuck in local minima**
- When **local minima exist**, Hill Climbing is **suboptimal**.

# Stochastic Hill Climbing


```
Input:  $Iter_{max}$ , ProblemSize
Output: Current
Current  $\leftarrow$  RandomSolution(ProblemSize)
For ( $iter_i \in Iter_{max}$ )
    Candidate  $\leftarrow$  RandomNeighbor(Current)
    If (Cost(Candidate)  $\geq$  Cost(Current))
        Current  $\leftarrow$  Candidate
    End
End
Return (Current)
```

- Random selection overcomes local minima
- The **selection probability** can vary with the steepness.
- Stochastic Hill Climbing is a **Local Optimization** algorithm
- Does **not require derivatives** of the search space.

# Simulated Annealing



```
S0 = GenerateSolution()
Temp = START_TEMP
K = 0 // iteration count
while(!stopping_condition(S0, Temp, K))
    S1 = Neighbor(S0) // make a neighbor of S0
    if(Fitness(S1) < Fitness(S0))
        S0 = S1 // if S1 is better than S0 take it
    else if(rand() < tempFunc(S0, S1, Temp, K))
        S0 = S1 // if S1 is worse than S0
                // conditionally take it
    end if
    AnnealingSchedule(S0, Temp, K) // anneal the temp
    K = K + 1 // increase the iteration count
end while
```



- **Simulated Annealing** = physics inspired twist on random walk
- Escape **local maxima** by allowing some bad moves
- Instead of picking the best move, **pick one randomly**

# Genetic Algorithm


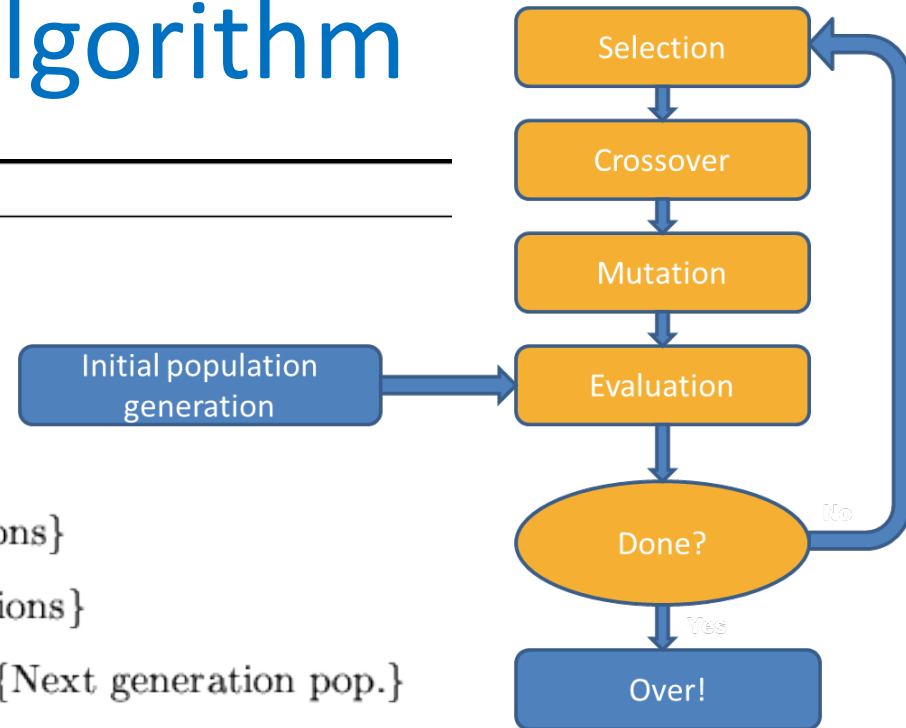
---

**Algorithm 1** Pseudocode for a Genetic Algorithm

---

```
1:  $t \leftarrow 0$ ;  
2: InitPopulation[ $P(t)$ ]; {Initializes the population}  
3: EvalPopulation[ $P(t)$ ]; {Evaluates the population}  
4: while not termination do  
5:    $P'(t) \leftarrow \text{Variation}[P(t)]$ ; {Creation of new solutions}  
6:   EvalPopulation[ $P'(t)$ ]; {Evaluates the new solutions}  
7:    $P(t+1) \leftarrow \text{ApplyGeneticOperators}[P'(t) \cup Q]$ ; {Next generation pop.}  
8:    $t \leftarrow t + 1$ ;  
9: end while
```

---

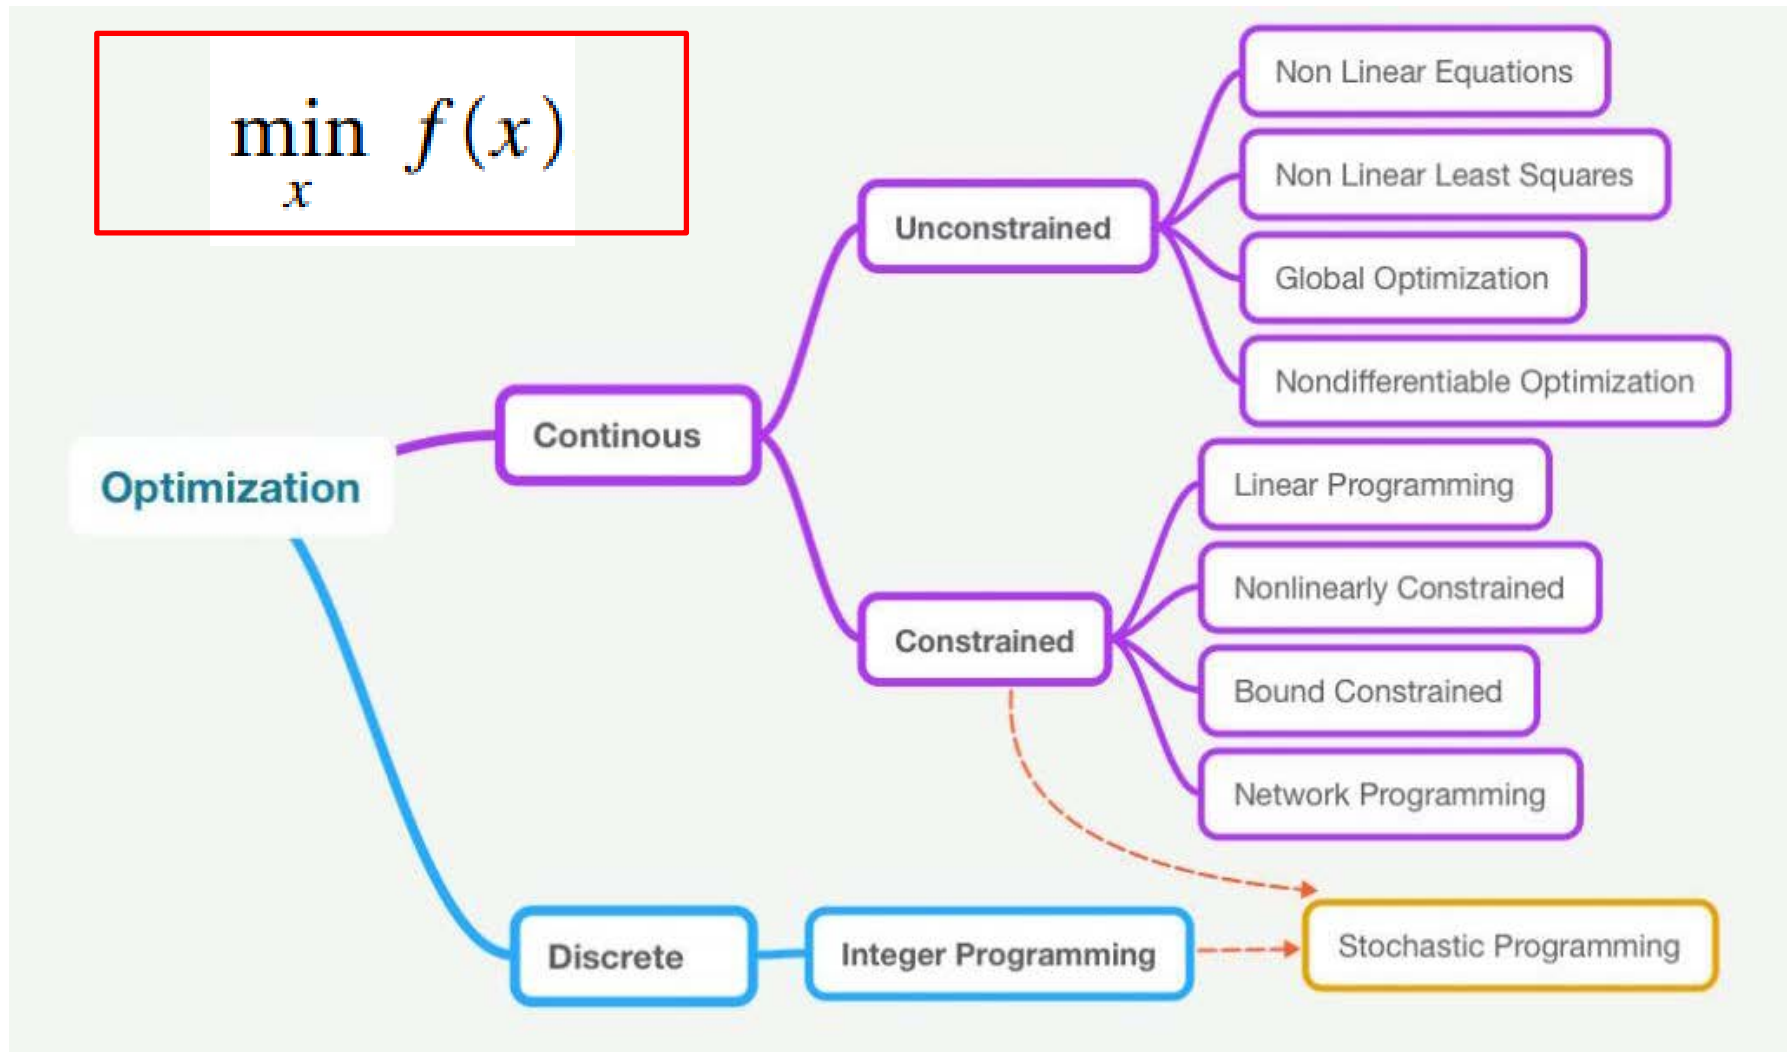


- **Twist on Local Search**
- **Successor** is generated by **combining parent states**
- A **state** is represented as a **string over an alphabet**
- **Randomly generated states** (population)



## Part V – Optimizing



# Numerical Optimization



# Optimization Approaches

- **Derivative free optimization** refers to problems for which derivative information is unavailable or impractical to obtain.
- **Gradient-based optimization** : method with the search directions defined by the *gradient* of the function at the current point.
- **Stochastic optimization** generate and use random variables or which involve random objective functions or random constraints.
- **Constrained optimization** optimize an objective function with respect to some variables in the presence of constraints.

# Numerical Optimization




- **Convex** : any local minimum is a global minimum
- **Non-Convex** : bad local minima, stay stuck in a local minima that is not a global minima.

# Derivatives and Gradient

- Central **Finite Differencing**


$$\frac{\partial f}{\partial x_i}(x) \approx \frac{f(x + \epsilon e_i) - f(x - \epsilon e_i)}{2\epsilon}.$$



- **Gradient** methods:



$$\nabla f = \left( \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$

to increase/reduce  $f$ , e.g., by  $\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$



- The gradient gives the steepest descent direction.
- Derivatives may not exist **or too costly to compute.**

# Rosenbrock Function



$$f(x, y) = (1 - x)^2 + 100(y - x^2)^2$$

*The **Rosenbrock function** is a **non-convex function**, introduced by Howard H. Rosenbrock in 1960, which is used as a performance test problem for **optimization algorithms**.*

# Gradient Descent

Performing a line search along the  
direction of the gradient

---

**Algorithm:** Gradient Descent

---

**Input:**  $Y, \Theta, \mathbf{X}, \alpha$ , tolerance, max iterations

**Output:**  $\Theta$

**for**  $i = 0; i < \text{max iterations}; i++$  **do**

    current cost =  $\text{Cost}(Y, \mathbf{X}, \Theta)$

**if**  $\text{current cost} < \text{tolerance}$  **then**


        | break

**else**

        | gradient =  $\text{Gradient}(Y, \mathbf{X}, \Theta)$

        |  $\theta_j \leftarrow \theta_j - \alpha \cdot \text{gradient}$

---





# Gradient Descent

---



## Algorithm 3.2 Gradient Descent

---

- 1: **Input:** Initial point  $w_0$ , gradient norm tolerance  $\epsilon$
  - 2: Set  $t = 0$
  - 3: **while**  $\|\nabla J(w_t)\| \geq \epsilon$  **do**
  - 4:    $w_{t+1} = w_t - \eta_t \nabla J(w_t)$
  - 5:    $t = t + 1$
  - 6: **end while**
  - 7: **Return:**  $w_t$
- 

- **Gradient** of the input function:  $\nabla J(w_t)$
- Decaying **Stepsize**:  $\eta_t = 1/\sqrt{t}$ .
- Gradient Descent **Update**:

$$w_{t+1} = w_t - \eta_t \nabla J(w_t)$$






# Stochastic Gradient Descent

---

**Algorithm 3.8** Stochastic Gradient Descent

---

- 1: **Input:** Maximum iterations  $T$ , batch size  $k$ , and  $\tau$
  - 2: Set  $t = 0$  and  $w_0 = 0$
  - 3: **while**  $t < T$  **do**
  - 4:   Choose a subset of  $k$  data points  $(x_i^t, y_i^t)$  and compute  $\nabla J_t(w_t)$
  - 5:   Compute stepsize  $\eta_t = \sqrt{\frac{\tau}{\tau+t}}$
  - 6:    $w_{t+1} = w_t - \eta_t \nabla J_t(w_t)$
  - 7:    $t = t + 1$
  - 8: **end while**
  - 9: **Return:**  $w_T$
- 



- **Very noisy estimation of the gradient**
- **Decaying Stepsize:**

$$\eta_t = \frac{\tau}{\tau + t} \quad \tau > 0$$


- **Gradient Descent Updates:**
- $$w_{t+1} = w_t - \eta_t \nabla J(w_t)$$

# Linear Least Squares

$$\mathbf{Ax} = \mathbf{b}$$

$$\begin{bmatrix} 1 & 2 & -2 & -6 \\ -3 & -1 & -2 & \alpha \\ -4 & 3 & 9 & 16 \\ 5 & 7 & -6 & -15 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

$\uparrow$                        $\uparrow$                        $\uparrow$   
 $\mathbf{A}$                        $\mathbf{x}$                        $\mathbf{b}$




$$f(\mathbf{x}) = \frac{1}{2} \|\mathbf{Ax} - \mathbf{b}\|^2$$

$$\min_x \|Ax - b\|_2$$




*Data fitting **estimated parameters** to optimize the fitting. The solution is obtained by solving the **normal equation**.*

$$\nabla f(x) = \mathbf{A}^T \mathbf{Ax} - \mathbf{A}^T \mathbf{b}$$

# Constrained Optimization



$$\begin{array}{ll} \min_{\mathbf{x}} & f(\mathbf{x}) \\ \text{subject to} & \mathbf{x} \in C, \end{array}$$



**Constrained optimization** is the process of optimizing an objective function in the presence of **constraints** on variables.

## Part VI – Conclusions

# Game vs Search

- **Search:**

- Solution is **heuristic** method for finding goal
- Find **(near-)optimal** solution
- **Evaluation** of a **cost function**

- **Game:**

- Solution is **strategy**.
- **Approximate** solution.
- **Unpredictable** opponent.
- Kind of **uncertainty**.



# Limitation of Game Theory

- Some assumptions are not practical.
- Increasing complexity with the increase of the players.
- **Infinite number of strategy.**
- **Outcome: the gain of one person is the loss of another person.**
- **Knowledge about strategy: players as the knowledge of strategies.**
- **Assumption of maximin and minimax**
- Assumption that players are equally wise and behave rationally.
- Risk and Certainty of Pay Offs / **Rules of game**

# Lab Activities

- **Activity 1:** Payoff Matrix and Equilibrium (30 min)
- **Activity 2:** Stochastic Hill Climbing (30 min)
- **Break** (no break)
- **Activity 3:** Backtracking and Search (30 min)
- **Activity 4:** Gradient Decent (30 min)

# References

- [1] - Problem Solving, search and Control strategy - myreaders.info
- [2] - Heuristic Search - Andrea Torsello
- [3] - Problem Solving and Search - Russell and Norvig
- [4] - Artificial Intelligence - Nancy E. Reed
- [5] - Artificial Intelligence - Tabu Search - Fred Glover
- [6] - Numerical Optimization Techniques - Leon Bottou
- [7] - Search Techniques for Artificial Intelligence - Ulle Endriss
- [8] - Artificial Intelligence Programming in Prolog - Tim Smith
- [9] - Recursive Backtracking - Marty Stepp and Helene Martin
- [10] - Computational Game Theory - University of Oxford
- [11] - Introduction to Artificial Intelligence State Space Search Gregory Adam
- [12] - AI Techniques for Game Programming - Team LRN
- [13] - Uninformed search methods - Doina Precup
- [14] - Algorithmic Game theory - Noam Nisan
- [15] - A tutorial on Tabu Search - Hertz et al
- [16] - Local search algorithms - Peter Ljunglof
- [17] - Problem Solving and Search in Artificial Intelligence Basic Concepts - Nysret Musliu
- [18] - Artificial Intelligence - Game Theory
- [19] - Solving Problems by Searching - Marco Chiarandini - Stuart Russell and Peter Norvig
- [20] - Graduate AI Computational Game Theory - Michael Benisch



# References

- [21] - Artificial Intelligence For Games - Ian Millington and John Funge
- [22] - Introduction to AI Techniques: Game Search, Minimax, and Alpha Beta Pruning
- [23] - Game Theory - University of Notre Dame
- [24] - Introduction to informed search The A\* search algorithm - Penn University
- [25] - Least Squares with Examples in Signal Processing Ivan Selesnick
- [26] - Linear Least Squares Optimization - Kurt Bryan
- [27] - Nonlinear Optimization Least Squares Problems - Nicolas Borlin
- [28] - Adversarial Search and GamePlaying - Oliver Schulte
- [29] - Numerical Optimization - Nocedal and Wright
- [30] - Introduction to non-linear optimization - Ross A. Lippert
- [31] - State Space Representation and Search - Prof. Hugh Murrell
- [32] - Fundamentals of Artificial Intelligence - Alan Smaill
- [33] - Iterated Local Search Variable Neighborhood Search - Thomas Stutzle
- [34] - Backtracking Search Algorithms - Peter van Beek
- [35] - Tabu Search - Fred Glover and Rafael Marti
- [36] - Based on slides of Padhraic Smyth, Stuart Russell, Rao Kambhampati, Raj Rao, Dan Weld...